

# Training neural network for Assignment 4, LED segments

## Table of Contents

### Training neural network for Assignment 4, LED segments

- Define a generic NN layer
- Define a network as a concatenation of many layers
- Create a two-layer network
- Create Nueral Net with 2 layers
- Train on Input data
- Check tolerance
- Plot the result
  - Accuracy over time
  - Loss Curve
  - Bar Plot interaction

```
1 begin
2     using Plots
3
4     # Packages for automatic differentiation and neural networks
5     using Flux, Zygote
6 end
```

## Define a generic NN layer

No changes here from the lecture notes

```
1 begin
2     struct Layer
3         W::Matrix{Float32} # weight matrix - Float32 for faster gradients
4         b::Vector{Float32} # bias vector
5         activation::Function
6         Layer(in::Int64,out::Int64,activation::Function=identityFunction) =
7             new(randn(out,in),randn(out),activation) # constructor
8     end
9
10    (m::Layer)(x) = m.activation.(m.W*x .+ m.b) # feed-forward pass
11 end
```

Define some required activation functions. ReLu is a standard neural network activation function.

```
1 begin
2     ReLu(x) = max(0,x)
3     identityFunction(x) = x
4 end;
```

## Define a network as a concatenation of many layers

No changes here from notes

```

1 begin
2     struct Network
3         layers::Vector{Layer}
4         Network(layers::Vararg{Layer}) = new(vcat(layers...))
5         # constructor - allow arbitrarily many layers
6     end
7
8     (n::Network)(x) = reduce((left,right)->right◦left, n.layers)(x)
9     # perform layer-wise operations over arbitrarily many layers
10 end

```

## Create a two-layer network

Input is numerical encoded segment values, see **inputs** matrice

Output is one-hot-encoded digits from 0-9, see **outputs** matrice

mse (generic function with 1 method)

```

1 begin
2     #inputs = collect(-3:0.1:3) # create training data
3     #targetOutput = sin.(inputs)
4
5     inputs = [
6         1 1 1 1 1 1 0; # 0
7         0 1 1 0 0 0 0; # 1
8         1 1 0 1 1 0 1; # 2
9         1 1 1 1 0 0 1; # 3
10        0 1 1 0 0 1 1; # 4
11        1 0 1 1 0 1 1; # 5
12        1 0 1 1 1 1 1; # 6
13        1 1 1 0 0 0 0; # 7
14        1 1 1 1 1 1 1; # 8
15        1 1 1 1 0 1 1; # 9
16    ]
17
18    targetOutput = [
19        1 0 0 0 0 0 0 0 0; # 0
20        0 1 0 0 0 0 0 0 0; # 1
21        0 0 1 0 0 0 0 0 0; # 2
22        0 0 0 1 0 0 0 0 0; # 3
23        0 0 0 0 1 0 0 0 0; # 4
24        0 0 0 0 0 1 0 0 0; # 5
25        0 0 0 0 0 0 1 0 0; # 6
26        0 0 0 0 0 0 0 1 0; # 7
27        0 0 0 0 0 0 0 0 1; # 8
28        0 0 0 0 0 0 0 0 1; # 9
29    ]
30    mse(x,y) = sum((x .- y).^2)/length(x) # MSE will be our loss function
31 end

```

## Create Nueral Net with 2 layers

Input layer is 7 neurons, first layer is 100 neurons

Second layer is 100 neurons and output layer is 10 neurons

```

Network([Layer(100x7 Matrix{Float32}:
      1  22684      -1  22644      -2  72411      -1  65359      -0  790929      1  12893      -1

```

```

1 begin
2     using Random
3     Random.seed!(54321) # for reproducibility
4
5     twoLayerNeuralNet = Network(Layer(7,100 ,ReLU), Layer(100,10))
6     # instantiate a two-layer network
7 end

```

## Train on Input data

---

- This involves 2 loops, one nested inside another
- The outer loop iterates through 'epochs'

The inner loop iterates 10 times per epoch, for each loop we iterate over 10 random digit patterns, so 'i' is a random value between 1-10 (not-recurring) for each loop. This is done to ensure the input digit (vector) is always random.

The outer loop runs 800 times, 800 epochs, inner loop 10 times per epoch

Have the following variables

- input - LED vector, the LED number
- output - One-hot-encoded digit
- prediction - 10 digit vector, each value is confidence the corresponding index is the predicted output for the LED input

Training remains same as in notes, adjusting parameters based on the gradients, calculated by differentiating the mean squared error with respect to the parameters

For graphing have the the following arrays

- epochCount - for reference y axis later
- epochAccuracy - store accuracy (correct predictions out of 10) for each epoch
- lossCurve - store mean squared error

Accuracy is calculated by checking if predicted digit is equal to target digit, if so increment a counter 'correct' (this is done inside the inner loop). The value of 'correct' is divided by 10 and pushed to the epochAccuracy array

```

1 begin
2
3     Flux.@functor Layer      # set the Layer-struct as being differentiable
4     Flux.@functor Network   # set the Network-struct as being differentiable
5
6     parameters = Flux.params(twoLayerNeuralNet)
7         # obtain the parameters of the layers (recurses through network)
8
9     optimizer = ADAM(0.001) # from Flux-library
10
11     epochCount = []
12     epochAccuracy = []
13     lossCurve = []
14     epochs = 800
15
16     for epoch in 1:epochs
17         correct = 0
18         loss_count = 0
19
20         for j in shuffle(1:10)
21             # Select input-output pair
22             input = inputs[j, :]
23             output = targetOutput[j, :]
24
25             #run the network based on the input
26             prediction = twoLayerNeuralNet(input)
27             loss = mse(prediction, output)
28
29             #calculate the gradient
30             gradients = Zygote.gradient(() -> mse(twoLayerNeuralNet(input), output),
31                 parameters)
32
33             #adjust the parameters based on the gradient
34             Flux.Optimise.update!(optimizer, parameters, gradients)
35
36
37             loss_count += loss
38
39
40             out_label = findmax(output)[2] - 1
41             pred_label = findmax(prediction)[2] - 1
42             if out_label == pred_label
43                 correct += 1
44             end
45         end
46
47
48         push!(epochAccuracy, correct / 10)
49         push!(lossCurve, loss_count / 10)
50         push!(epochCount, epoch)
51     end
52
53
54 end

```

## Check tolerance

To ensure the elements of the output vector can be within a 0.1 tolerance of the target value I use monte carlo method.

A random digit from 0-9 is inputted into the model to get a prediction vector, the prediction vector values are checked to ensure they are within the tolerance. This is done 800 times

- Predicting confidence - highest value in vector, what network predicts digit is
- Non-predicting confidence - all other values in vector beside predicting confidence

Over 800 epochs, the highest (max) and lowest (min) values are recorded for the

- predicting confidence - should be between 0.9 and 1.1
- Non-predicting confidence - should be between -0.1 and 0.1

checkNonPrediction (generic function with 1 method)

```

1 function checkNonPrediction(prediction, pred_value, max_nonPred_confidence,
  min_nonPred_confidence)
2     for i in 1:10
3         if i != pred_value+1
4             if prediction[i] > max_nonPred_confidence
5                 max_nonPred_confidence = prediction[i]
6             elseif prediction[i] < min_nonPred_confidence
7                 min_nonPred_confidence = prediction[i]
8             end
9         end
10    end
11    return (max_nonPred_confidence, min_nonPred_confidence)
12 end

```

```

1 begin
2     max_pred_confidence = 0
3     min_pred_confidence = 100
4
5     max_nonPred_confidence = 0
6     min_nonPred_confidence = 100
7     for i in 1:500
8         pred_confidence_count = 0
9         for j in shuffle(1:10)
10            input = inputs[j, :]
11            output = targetOutput[j, :]
12            prediction = twoLayerNeuralNet(input)
13
14            input_label = join(inputs[j, :], ", ")
15            out_value = findmax(output)[2] - 1
16            pred_value = findmax(prediction)[2] - 1
17
18            pred_confidence = findmax(prediction)[1]
19
20            #Check if confidence is between 0.1 tolerance, so between 0.9 and 1.1
21            if pred_confidence > max_pred_confidence
22                max_pred_confidence = pred_confidence
23            elseif(pred_confidence < min_pred_confidence)
24                min_pred_confidence = pred_confidence
25            end
26            tuple = checkNonPrediction(prediction, pred_value,
27                                     max_nonPred_confidence,min_nonPred_confidence )
28
29            max_nonPred_confidence = tuple[1]
30            min_nonPred_confidence = tuple[2]
31
32        end
33        pred_confidence_count = (pred_confidence_count+1)
34    end
35    println("Max prediction confidence: $max_pred_confidence\nMin prediction
36    confidence: $min_pred_confidence\n")
37
38    println("Max Non Prediction confidence: $max_nonPred_confidence\nMin Non
39    Prediction confidence: $min_nonPred_confidence\n")
40
41    println("Is value prediction outside tolerance: ", (max_pred_confidence > 1.1 ||
42    min_pred_confidence < 0.9))
43    println("Is non-value prediction outside tolerance: ", (max_nonPred_confidence >
44    0.1 || min_nonPred_confidence < -0.1))
45 end

```

```

Max prediction confidence: 1.006978
Min prediction confidence: 0.9844361

Max Non Prediction confidence: 0.032113
Min Non Prediction confidence: -0.046803057

Is value prediction outside tolerance: false
Is non-value prediction outside tolerance: false

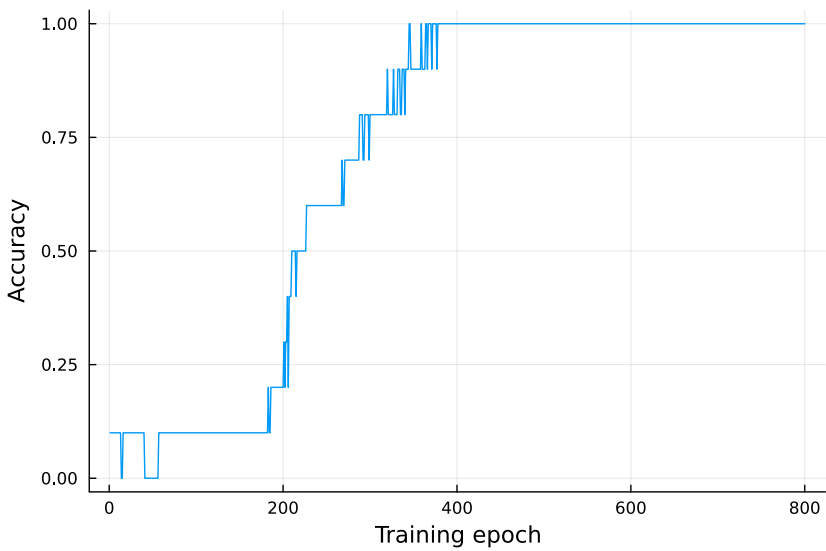
```

# Plot the result

## Accuracy over time

- X-axis - epochs, number of times trained model 10 times on random inputs 1-10, 800 times
- Y-axis - Accuracy, from 0-1

See how accuracy is stable at 1.0 at 400 epochs, around halfway, I think this ensures not under/over-confident predictions, hence stay within tolerance




```
1 begin
2   plot(epochCount, epochAccuracy, lab="")
3   yaxis!("Accuracy")
4   xaxis!("Training epoch")
5 end
```

## Loss Curve

- X-axis - epochs
- Y-axis - Mean Squared Error


lossCurvePlot =

Loss curve



```
1 lossCurvePlot = plot(lossCurve, title = "Loss curve", legend=:none)
```

## Bar Plot interaction



```

1 begin
2   function get_prediction(digit)
3     input = inputs[digit + 1, :] # +1 why?? julia -1 indexed??
4     return twoLayerNeuralNet(input)
5   end
6
7   function plot_prediction(digit)
8     prediction = get_prediction(digit)
9
10    input_str = join(inputs[digit + 1, :], ", ")
11    target_output = join(targetOutput[digit + 1, :], ", ")
12    predicted_output = join(targetOutput[findmax(prediction)[2], :], ", ")
13
14    Plots.bar(0:9, prediction, title="Prediction for LED $digit", legend=false,
15    xticks=0:9)
16    plot!(ylim=(0, 1.2))
17    annotate!(2, 1.2, text("Input: [$input_str]", 10))
18    annotate!(3.2, 1.15, text("Target Output: [$target_output]", 10))
19    annotate!(3.5, 1.1, text("Predicted Output: [$predicted_output]", 10))
20    #why does x shift? 2 then 3.2 then 3.5???
21
22  end
23
24  plot_prediction(digit)
25 end

```

8

```
1 @bind digit Slider(0:9, show_value=true)
```