

Table of Contents

Training a CNN classifier using Flux
Loading the dataset
Training the network
Testing the network

```
1 begin
2     using PlutoUI
3     using Latexify
4     TableOfContents()
5 end
```

Training a CNN classifier using Flux

MNIST is a dataset of 60k small training images of handwritten digits from 0 to 9.

We will do the following steps in order:

- Load MNIST training and test datasets
- Define a Convolution Neural Network (CNN)
- Define a loss function
- Train the network on the training data
- Test the network on the test data

Loading the dataset

Metalhead.jl is an excellent package that has a number of predefined and pretrained computer vision models. It also has a number of dataloaders that come in handy to load datasets.

```
1 begin
2     using Statistics
3     using CUDA
4     using Flux, Flux.Optimise
5     using MLDatasets: CIFAR10
6     using Images.ImageCore
7     using Images
8     using Flux: onehotbatch, onecold
9     using Base.Iterators: partition
10
11     using Plots
12 end
```

Package cuDNN not found in current path.
- Run `import Pkg; Pkg.add("cuDNN")` to install the cuDNN package, then restart Julia.
- If cuDNN is not installed, some Flux functionalities will not be available when running on the GPU.

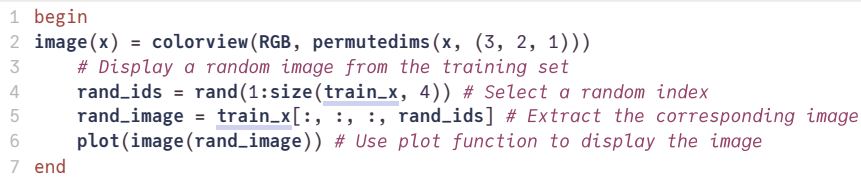
Set an environment variable to stop the system asking for approval to download data.

```
1 ENV["DATADEPS_ALWAYS_ACCEPT"] = true;
```

```
1 begin
2     train_x, train_y = CIFAR10(split=:train)[:]; #loadx is image data, y is label
3     train_x = reshape(train_x, 32, 32, 3, :) #turn x array images to 3d array
4     labels2 = onehotbatch(train_y, 0:9) #convert training labels to one hot encoded
5     classes = ["airplane", "automobile", "bird", "cat", "deer", "dog", "frog",
6               "horse", "ship", "truck"] #label array
7 end;
```

Tip

The size and dimensionality of the CIFAR10 data base are different: there are 50,000 images of 32 x 32 pixels involving 3 colour channels



Use this function: `image(x) = colorview{RGB, permutedims(x, (3, 2, 1))}` to display the colour images instead of Gray.

The first 59k images (in batches of 1,000) will be our training set, and the rest are for validation used to track training progress. `partition` handily breaks down the set we give it in consecutive parts (1,000 in this case).

```

1 begin
2     train_x_small = train_x[:, :, :, 1:50000]
3     train_y_small = train_y[1:50000]
4     #labels2 = onehotbatch(train_y_small, 0:9)
5
6     train = ([ (train_x_small[:, :, :, i], labels2[:, i]) for i in partition(1:49000,
7     1000)] ) |> gpu
8
9     valset = 49001:50000
10    valX = train_x_small[:, :, :, valset] |> gpu #image data for validation set
11    valY = labels2[:, valset] |> gpu #label data for validation set
12 end

```

2/7

Tip

Because CIFAR10 comprises 50,000 images rather than 60,000, the split between training and validation needs to reflect this.

```
m3 = Chain(
  Conv((5, 5), 3 => 16, relu, pad=2), # 1_216 parameters
  MaxPool((2, 2)),
  Conv((5, 5), 16 => 8, relu, pad=2), # 3_208 parameters
  MaxPool((2, 2)),
  Main.var"#5#6"{typeof(reshape), typeof(size), Colon}(reshape, size, Colon()),
  Dense(512 => 120), # 61_560 parameters
  Dense(120 => 84), # 10_164 parameters
  Dense(84 => 10), # 850 parameters
  NNLib.softmax,
) # Total: 10 arrays, 76_998 parameters, 302.008 KiB.
```

```
1 #CNN work on batches in parrallell, so 1000 images go through it simulatneously
2 #change first conv layer 1 to 3, since rgb no longer greyscale
3 #input 32*32*3
4 m3 = Chain(
5   Conv((5,5), 3=>16, pad=(2,2), relu), #output remain 32*32, 16 channels now 16 not 3
6   MaxPool((2,2)), #reduce spatial dimension by half, 16*16 now
7   Conv((5,5), 16=>8, pad=(2,2), relu), #pad is 2 so spatial 16*16, channel is 8
8   MaxPool((2,2)), #half spatial dimension, so 8*8
9   #Now image dimension is 8*8*8 = 512
10  x -> reshape(x, :, size(x, 4)),
11  Dense(512, 120),
12  Dense(120, 84),
13  Dense(84, 10),
14  softmax) |> gpu
15
16 #First dense layer changed to take output of previous conv and pool layer
17 #First conv make is 16*16
18 #Second conv make 8*8
19 #La
```

Tip

You need to modify the input filter number of the first Conv layer and the input dimension of the first Dense layer.

We will use a crossentropy loss and the Momentum optimiser here. Crossentropy will be a good option when it comes to working with multiple independent classes. Momentum smooths out the noisy gradients and helps towards a smooth convergence. Gradually lowering the learning rate along with momentum helps to maintain a bit of adaptivity in our optimisation, preventing us from overshooting our desired destination.

```
1 using Flux: crossentropy, Momentum
```

```
1 begin
2   loss2(x, y) = sum(crossentropy(m3(x), y))
3   opt2 = Momentum(0.01)
4 end;
5
6 #m = 0.0001 - accuracy went down from 0.6950 to 0.6935
```

We can start writing our train loop where we will keep track of some basic accuracy numbers about our model. We can define an accuracy function for it like so:

accuracy (generic function with 1 method)

```
1 #compare predicted label to true label, onecold convert model output to class label
2 accuracy(x, y) = mean(onecold(m3(x), 0:9) .== onecold(y, 0:9))
```

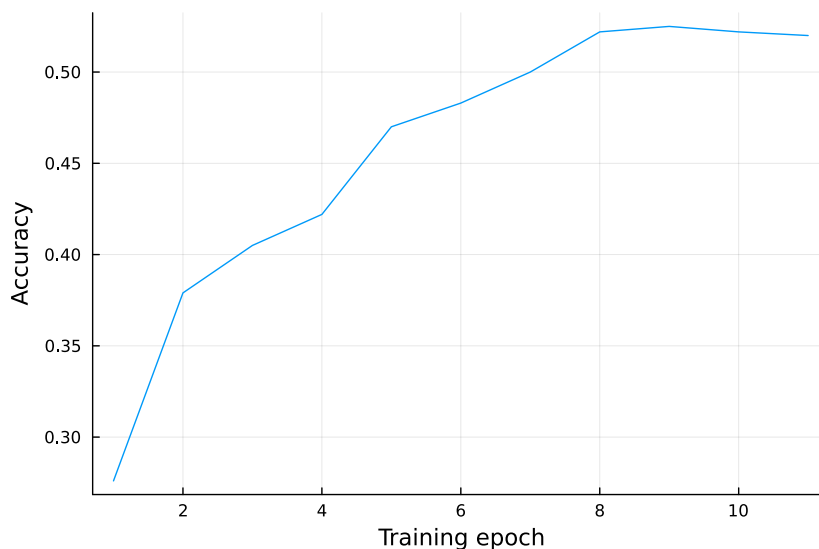
Training the network

Training is where we do a bunch of the interesting operations we defined earlier, and see what our net is capable of. We will loop over the dataset 10 times and feed the inputs to the neural network and optimise.

```

1 begin
2   train_acc = Float32[]
3   train_epochs = Int32[]
4   epochs = 11
5   for epoch = 1:epochs
6     #loop through each batch, so 59 iterations
7     #each batch, compute loss
8     for d in train
9       #gradient of loss function with respect to parms
10      # 'do' gradient() uses, so gradient for parms with ('do') loss
11      gs = gradient(Flux.params(m3)) do
12        l = loss2(d...) #loss function of current batch 'd', parallelized
13        #line above put batch through CNN, then compute loss
14      end
15      update!(opt2, Flux.params(m3), gs)
16    end
17    push!(train_acc, accuracy(valX, valY))
18    push!(train_epochs, epoch)
19  end
20 end

```



```

1 begin
2   plot(train_epochs, train_acc, lab="")
3   yaxis!("Accuracy")
4   xaxis!("Training epoch")
5 end

```

Tip

The training regime doesn't need modification.

Testing the network

We have trained the network for 10 passes over the training dataset. But we need to check if the network has learnt anything at all.

We will check this by predicting the class label that the neural network outputs, and checking it against the ground-truth. If the prediction is correct, we add the sample to the list of correct predictions. This will be done on a yet unseen section of data.

First step. Let us perform the exact same preprocessing on this set, as we did on our training set.

Next, display images from the test set.

```

1 begin
2   #load cifar10 test set
3   test_x, test_y = CIFAR10(split=:test)[:,:] #load x is image data, y is label
4   test_x = reshape(test_x, 32, 32, 3, :) #turn x array images to 3d array
5 end;

```

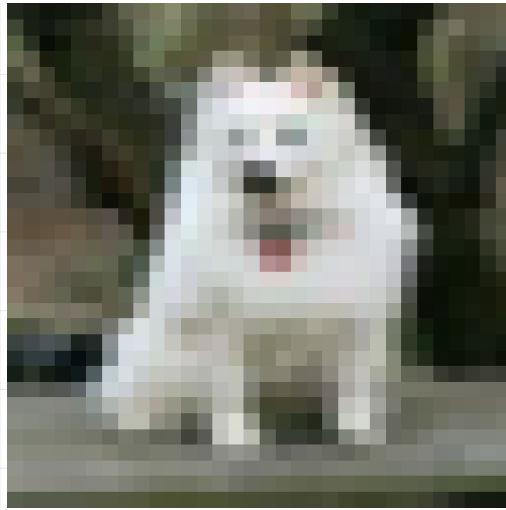
```

1 test_labels = onehotbatch(test_y, 0:9); #one hot encode the labels

```

Created test sets, for x and y. The code snippet below just displays the images

dog



```

1 begin
2   rnd_test = rand(1:size(test_x)[4]) #pick random indexes from total data set
3   plot(image(test_x[:, :, rnd_test]), axis=false) #plot image, no axis
4   annotate!(-1.0, 1.0, text(classes[test_y[rnd_test]+1], :blue, :right, 12))
   #annotate
5 end

```

The outputs are energies for the 10 classes. Higher the energy for a class, the more the network thinks that the image is of the particular class. Every column corresponds to the output of one image, with the 10 floats in the column being the energies.

Let's see how the model fared:

```

10x20 Matrix{Float32}:
0.0922064  0.0102547  0.0169114  ...  0.00156501  0.00903615  0.0987136
0.0363254  0.00481658  0.00421993  ...  0.00067695  0.00117676  0.732134
0.0317206  0.193057   0.0646506  ...  0.00112666  0.259643   0.0489655
0.0107373  0.201034    0.0758947  ...  0.37967    0.0940441  0.010446
0.043479   0.0705209   0.223071   ...  0.00164087  0.393149   0.0243906
0.00262163 0.152112    0.0614509  ...  0.601731   0.0501104  0.00513173
0.0227949  0.193075    0.241543   ...  0.0100709  0.188419   0.000536064
0.0058762  0.053236    0.283027   ...  0.0029483  0.00279208 0.000437009
0.561102   0.0670079   0.00330761 ...  0.000152686 0.000677557 0.00854503
0.193136   0.0548851   0.0259244  ...  0.000417025 0.000952252 0.0707004

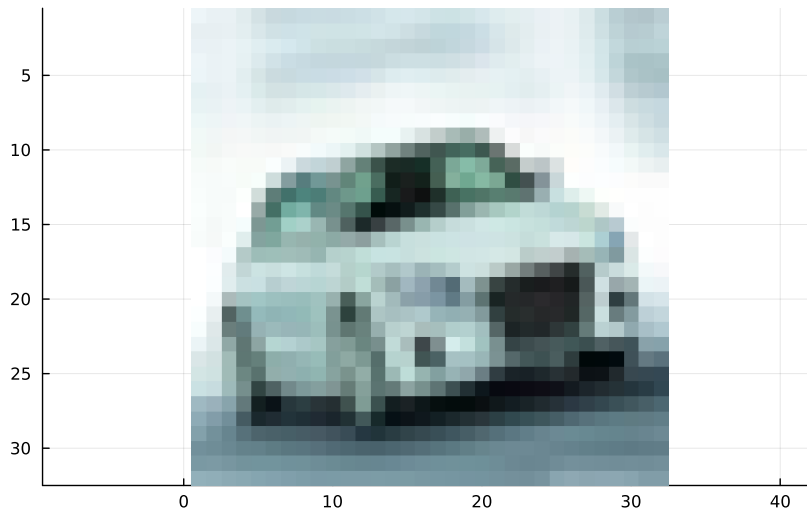
```

```

1 begin
2   ids = rand(1:size(test_x)[4], 20) #20 random ids
3   rand_test = test_x[:, :, ids] #get the images for the ids
4   rand_truth = test_y[ids] #true labels for selected ids
5   rand_out = m3(rand_test) #input random test images into neureal network
6 end

```

Predicted Label: automobile



```

1 begin
2     displayed_image = rand_test[:, :, :, inx] #get image at index of slider value
3     displayed_label_index = argmax(rand_out[:, inx]) # Get the index of the predicted
    label
4     predicted_label = classes[displayed_label_index]
5
6
7     plot(image(displayed_image)) # Display image
8     title!("Predicted Label: $predicted_label") #Display predicted label
9 end

```



```
1 @bind inx Slider(1:1:20, default=1)
```

Tip

For visualisation you should display performance on a random sample of test images, say 20, and set up a slider to navigate them while displaying the image, the ground truth label, and the label predicted by the network.

This looks similar to how we would expect the results to be. At this point, it's a good idea to see how our net actually performs on new data, that we have prepared.

0.538

```

1 begin
2
3     test = [(test_x[:, :, :, i], test_labels[:, i]) for i in partition(1:size(test_x,
4     4), 1000)]
4     accuracy(test[1]...)
5 end

```

This is much better than random chance set at 10% (since we only have 10 classes), and not bad at all for a relatively small hand-coded network like this one.

Let's take a look at how the net performed on all the classes performed individually.

```

1 begin
2     class_correct = zeros(10)
3     class_total = zeros(10)
4     for i in 1:10
5         preds = m3(test[i][1]) #makes prediction on first elem of the ith batch
6         lab = test[i][2] #Retrieves the true labels for the images in the batch
7         for j = 1:1000 #for each element in current batch
8             pred_class = findmax(preds[:, j])[2] #most predicted class for image
9             actual_class = findmax(lab[:, j])[2] #actual class for image
10            if pred_class == actual_class
11                class_correct[pred_class] += 1 #correct prediction increment counter
12            end
13            class_total[actual_class] += 1 #total count
14        end
15    end
16 end

```

	accuracy	class
1	0.517	"airplane"
2	0.614	"automobile"
3	0.375	"bird"
4	0.302	"cat"
5	0.462	"deer"
6	0.512	"dog"
7	0.729	"frog"
8	0.468	"horse"
9	0.76	"ship"
10	0.5	"truck"

```
1 begin
2     using DataFrames
3     DataFrame(accuracy=(class_correct ./ class_total), class=classes)
4 end
```

Tip
For legibility you should assign labels to the image categories.

The spread seems pretty good, with certain classes performing significantly better than the others.